



# Lecture 23: Course Overview

**Applied Machine Learning**

Volodymyr Kuleshov  
Cornell Tech

# **Congratulations on Finishing Applied Machine Learning!**

You have made it! This is our last machine learning lecture, in which we will do an overview of the different algorithms seen in the course.

# A Map of Applied Machine Learning

We will go through the following map of algorithms from the course.

## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] ● +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] ●

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] ● +

**Neural:** Perceptron [L14]   Multi-Layer Perceptron [L14] ●   ConvNet [L15] ●

**Tree-Based:** Decision Tree [L12] ●   Random Forest [L12] ●   (Grad) Boosting [L13] ●

Nearest Neighbors [L6] ●

## Unsupervised Learning

### Clustering

K-Means [L6]

Gaussian Mixture Models [L6] ▲

### Density Estimation

Histogram Method [L6]

Kernel Density Estimation [L6] +

### Dimensionality Reduction

Principal Component Analysis [L6]

● Non-Linear

▲ Probabilistic

+ Kernelized

■ Non-Parametric

■ Generative

# Supervised Machine Learning

At a high level, a supervised machine learning problem has the following structure:

$$\underbrace{\text{Dataset}}_{\text{Features, Attributes}} + \underbrace{\text{Learning Algorithm}}_{\text{Model Class} + \text{Objective} + \text{Optimizer}} \rightarrow \text{Predictive Model}$$

The predictive model is chosen to model the relationship between inputs and targets. For instance, it can predict future targets.

# Linear Regression

In linear regression, we fit a model

$$f_{\theta}(x) := \theta^{\top} \phi(x)$$

that is linear in  $\theta$ .

The features  $\phi(x) : \mathbb{R} \rightarrow \mathbb{R}^p$  are non-linear may non-linear in  $x$  (e.g., polynomial features), allowing us to fit complex functions.

## Supervised Learning

### Classification

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] ●

● Non-Linear

# Overfitting

Overfitting is one of the most common failure modes of machine learning.

- A very expressive model (a high degree polynomial) fits the training dataset perfectly.
- The model also makes wildly incorrect prediction outside this dataset, and doesn't generalize.

# Regularization

The idea of regularization is to penalize complex models that may overfit the data.

Regularized least squares optimizes the following objective (**Ridge**).

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n \left( y^{(i)} - \theta^\top \phi(x^{(i)}) \right)^2 + \frac{\lambda}{2} \cdot \|\theta\|_2^2.$$

If we use the L1 norm, we have the **LASSO**.



## Supervised Learning

### Classification

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

• Non-Linear

# Regression vs. Classification

Consider a training dataset  $\mathcal{D} = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})\}$ .

We distinguish between two types of supervised learning problems depending on the targets  $y^{(i)}$ .

1. **Regression:** The target variable  $y \in \mathcal{Y}$  is continuous:  $\mathcal{Y} \subseteq \mathbb{R}$ .
2. **Classification:** The target variable  $y$  is discrete and takes on one of  $K$  possible values:  $\mathcal{Y} = \{y_1, y_2, \dots, y_K\}$ . Each discrete value corresponds to a *class* that we want to predict.

## Supervised Learning

### Classification

Logistic Regression [L6]

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Nearest Neighbors [L6] •

• Non-Linear

■ Non-Parametric

# Parametric vs. Non-Parametric Models

Nearest neighbors is an example of a *non-parametric* model.

- A parametric model  $f_{\theta}(x) : \mathcal{X} \times \Theta \rightarrow \mathcal{Y}$  is defined by a finite set of parameters  $\theta \in \Theta$  whose dimensionality is constant with respect to the dataset
- In a non-parametric model, the function  $f$  uses the entire training dataset to make predictions, and the complexity of the model increases with dataset size.

- Non-parametric models have the advantage of not losing any information at training time.
- However, they are also computationally less tractable and may easily overfit the training set.

# Probabilistic vs. Non-Probabilistic Models

A probabilistic model is a probability distribution

$$P(x, y) : \mathcal{X} \times \mathcal{Y} \rightarrow [0, 1].$$

This model can approximate the data distribution  $P_{\text{data}}(x, y)$ .

If we know  $P(x, y)$ , we can use the conditional  $P(y|x)$  for prediction.

# Maximum Likelihood Learning

Maximum likelihood is an objective that can be used to fit any probabilistic model:

$$\theta_{\text{MLE}} = \arg \max_{\theta} \mathbb{E}_{x,y \sim \mathbb{P}_{\text{data}}} \log P(x, y; \theta).$$

It minimizes the KL divergence between the model and data distributions:

$$\theta_{\text{MLE}} = \arg \min_{\theta} \text{KL}(P_{\text{data}} \parallel P_{\theta}).$$

# Discriminative vs. Generative Models

There are two types of probabilistic models: *generative* and *discriminative*.

$$\underbrace{P_{\theta}(x, y) : \mathcal{X} \times \mathcal{Y} \rightarrow [0, 1]}_{\text{generative model}} \quad \underbrace{P_{\theta}(y|x) : \mathcal{X} \times \mathcal{Y} \rightarrow [0, 1]}_{\text{discriminative model}}$$

We can obtain predictions from generative models via  $\max_y P_{\theta}(x, y)$ .



## Supervised Learning

### Classification

Logistic Regression [L6]

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] ●

Ridge Regression [L4]

LASSO [L4]

Nearest Neighbors [L6] ●

● Non-Linear

▲ Probabilistic

■ Non-Parametric

■ Generative

# The Max-Margin Principle

Intuitively, we want to select linear decision boundaries with high *margin*.

This means that we are as confident as possible for every point and we are as far as possible from the decision boundary.

```
In [2]: import numpy as np
import pandas as pd
from sklearn import datasets

# Load the Iris dataset
iris = datasets.load_iris(as_frame=True)
iris_X, iris_y = iris.data, iris.target

# subsample to a third of the data points
iris_X = iris_X.loc[::4]
iris_y = iris_y.loc[::4]

# create a binary classification dataset with labels +/- 1
iris_y2 = iris_y.copy()
iris_y2[iris_y2==2] = 1
iris_y2[iris_y2==0] = -1

# print part of the dataset
pd.concat([iris_X, iris_y2], axis=1).head()
```

Out[2]:


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	-1
4	5.0	3.6	1.4	0.2	-1
8	4.4	2.9	1.4	0.2	-1
12	4.8	3.0	1.4	0.1	-1
16	5.4	3.9	1.3	0.4	-1

```
In [3]: # https://scikit-learn.org/stable/auto\_examples/neighbors/plot\_classification.html
%matplotlib inline
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [12, 4]
import warnings
warnings.filterwarnings("ignore")

# create 2d version of dataset and subsample it
X = iris_X.to_numpy()[::2]
x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
xx, yy = np.meshgrid(np.arange(x_min, x_max, .02), np.arange(y_min, y_max, .02))

# Plot also the training points
p1 = plt.scatter(X[:, 0], X[:, 1], c=iris_y2, s=60, cmap=plt.cm.Paired)
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
plt.legend(handles=p1.legend_elements()[0], labels=['Setosa', 'Not Setosa'], loc=
'lower right')
```

```
Out[3]: <matplotlib.legend.Legend at 0x12b41fb00>
```



```

In [4]: from sklearn.linear_model import Perceptron, RidgeClassifier
        from sklearn.svm import SVC
        models = [SVC(kernel='linear', C=10000), Perceptron(), RidgeClassifier()]


        def fit_and_create_boundary(model):
            model.fit(X, iris_y2)
            Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
            Z = Z.reshape(xx.shape)
            return Z




        plt.figure(figsize=(12,3))
        for i, model in enumerate(models):
            plt.subplot('13%d' % (i+1))
            Z = fit_and_create_boundary(model)
            plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)

            # Plot also the training points
            plt.scatter(X[:, 0], X[:, 1], c=iris_y2, edgecolors='k', cmap=plt.cm.Paired)
            if i == 0:
                plt.title('Good Margin')
            else:
                plt.title('Bad Margin')
            plt.xlabel('Sepal length')
            plt.ylabel('Sepal width')

        plt.show()

```





## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] ●

Ridge Regression [L4]

LASSO [L4]

Nearest Neighbors [L6] ●

● Non-Linear

▲ Probabilistic

■ Non-Parametric

■ Generative



# The Kernel Trick

Many algorithms in machine learning only involve dot products  $\phi(x)^\top \phi(z)$  but not the features  $\phi$  themselves.

We can often compute  $\phi(x)^\top \phi(z)$  very efficiently for complex  $\phi$  using a kernel function  $K(x, z) = \phi(x)^\top \phi(z)$ . This is the **kernel trick**.

## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] • +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] • +

Nearest Neighbors [L6] •

• Non-Linear

▲ Probabilistic

+ Kernelized

Non-Parametric

Generative

# Tree-Based Models

Decision trees output target based on a tree of human-interpretable decision rules.

- **Random forests** combine large trees using *bagging* to reduce overfitting.
- **Boosted trees** combine small trees to reduce underfitting.

## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] • +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] • +

**Tree-Based:** Decision Tree [L12] • Random Forest [L12] • (Grad) Boosting [L13] •

Nearest Neighbors [L6] •

• Non-Linear

▲ Probabilistic

+ Kernelized

Non-Parametric

Generative

# Neural Networks

Neural network models are inspired by the brain.

- A Perceptron is an artificial model of a neuron.
- MLP stack multiple layers of artificial neurons.
- ConvNets tie the weights of neighboring neurons into receptive fields that implement the convolution operation.

## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] • +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] • +

**Neural:** Perceptron [L14]   Multi-Layer Perceptron [L14] •   ConvNet [L15] •

**Tree-Based:** Decision Tree [L12] •   Random Forest [L12] •   (Grad) Boosting [L13] •

Nearest Neighbors [L6] •

• Non-Linear

▲ Probabilistic

+ Kernelized

Non-Parametric

Generative

# Unsupervised Learning

We have a dataset *without* labels. Our goal is to learn something interesting about the structure of the data:

- **Clusters** hidden in the dataset.
- A **low-dimensional representation** of the data.
- Recover the **probability density** that generated the data.

## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] • +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] • +

**Neural:** Perceptron [L14]   Multi-Layer Perceptron [L14] •   ConvNet [L15] •

**Tree-Based:** Decision Tree [L12] •   Random Forest [L12] •   (Grad) Boosting [L13] •

Nearest Neighbors [L6] •

## Unsupervised Learning

### Clustering

K-Means [L6]

Gaussian Mixture Models [L6] ▲

• Non-Linear

▲ Probabilistic

+ Kernelized

■ Non-Parametric

■ Generative



## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] • +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] • +

**Neural:** Perceptron [L14] Multi-Layer Perceptron [L14] • ConvNet [L15] •

**Tree-Based:** Decision Tree [L12] • Random Forest [L12] • (Grad) Boosting [L13] •

Nearest Neighbors [L6] •

## Unsupervised Learning

### Clustering

K-Means [L6]

Gaussian Mixture Models [L6] ▲

### Density Estimation

Histogram Method [L6]

Kernel Density Estimation [L6] +

• Non-Linear

▲ Probabilistic

+ Kernelized

Non-Parametric

Generative

## Supervised Learning

### Classification

Logistic Regression [L6]

Support Vector Machine [L9, L10]

Kernelized SVM [L11] • +

Gaussian Discr. Analysis [L7] ▲

Naïve Bayes [L8] ▲

### Regression

Ordinary Least Squares [L3]

Non-Linear Least Squares [L3] •

Ridge Regression [L4]

LASSO [L4]

Kernelized Ridge [L11] • +

**Neural:** Perceptron [L14] Multi-Layer Perceptron [L14] • ConvNet [L15] •

**Tree-Based:** Decision Tree [L12] • Random Forest [L12] • (Grad) Boosting [L13] •

Nearest Neighbors [L6] •

## Unsupervised Learning

### Clustering

K-Means [L6]

Gaussian Mixture Models [L6] ▲

### Density Estimation

Histogram Method [L6]

Kernel Density Estimation [L6] +

### Dimensionality Reduction

Principal Component Analysis [L6]

• Non-Linear

▲ Probabilistic

+ Kernelized

Non-Parametric

Generative

# How To Decide Which Algorithm to Use

One factor is how much data you have. In the **small data** (<10,000) regime, consider:

- Linear models with hand-crafted features (LASSO, LR, NB, SVMs)
- Kernel methods often work best (e.g., SVM + RBF kernel)
- Non-parametric methods (kernels, nearest neighbors) are also powerful

In the **big data** regime,

- If using "high-level" features, gradient boosted trees are state-of-the-art
- When using "low-level" representations (images, sound signals), neural networks work best
- Linear models with good features are also good and reliable

Some additional advice:

- If interpretability matters, use decision trees or LASSO.
- When uncertainty estimates are important use probabilistic methods.
- If you know the data generating process, use generative models.

# What's Next? Ideas for Courses

Consider the following courses to keep learning about ML:

- Graduate courses in the Spring semester at Cornell (generative models, NLP, etc.)
- Masters courses: Deep Learning Clinic, ML Engineering, Data Science, etc.
- Online courses, e.g. Full Stack Deep Learning

# What's Next? Ideas for Research

In order to get involved in research, I recommend:

- Contacting research groups at Cornell for openings
- Watching online ML tutorials, e.g. NeurIPS
- Reading and implementing ML papers on your own


# What's Next? Ideas for Industry Projects

Finally, a few ideas for how to get more practice applying ML in the real world:

- Participate in Kaggle competitions and review solutions
- Build an open-source project that you like and host it on Github



# Thank You For Taking Applied Machine Learning!



In [ ]: